

Enron Submission Free-Response Questions

A critical part of machine learning is making sense of your analysis process and communicating it to others. The questions below will help us understand your decision-making process and allow us to give feedback on your project. Please answer each question; your answers should be about 1-2 paragraphs per question. If you find yourself writing much more than that, take a step back and see if you can simplify your response!

When your evaluator looks at your responses, he or she will use a specific list of rubric items to assess your answers. Here is the link to that rubric: [\[Link\]](#) Each question has one or more specific rubric items associated with it, so before you submit an answer, take a look at that part of the rubric. If your response does not meet expectations for all rubric points, you will be asked to revise and resubmit your project. Make sure that your responses are detailed enough that the evaluator will be able to understand the steps you took and your thought processes as you went through the data analysis.

Once you've submitted your responses, your coach will take a look and may ask a few more focused follow-up questions on one or more of your answers.

We can't wait to see what you've put together for this project!

1. Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those? [relevant rubric items: "data exploration", "outlier investigation"]

This data set includes a list of Enron employees. The 'poi' column is a binary column that states whether they were a person of interest in the Enron case. The rest of the columns are various statistics about each person, and we will be using those statistics to run machine learning algorithms and see if we can get one that can make decent predictions about who is a poi or not.

There were some outliers in the data, and we used an outlier cleaner to cap the top and bottom of the data using a z-score method. Anything above a z-score of 3 was capped at a z-score of 3, and those data points are printed to the console.

```
[5 rows x 20 columns]
Cleaning Outliers
Outliers capped for feature salary for person Index(['TOTAL'], dtype='object')
Outliers capped for feature to_messages for person Index(['KITCHEN LOUISE', 'SHAPIRO RICHARD S', 'KEAN STEVEN J'], dtype='object')
Outliers capped for feature deferral_payments for person Index(['TOTAL'], dtype='object')
Outliers capped for feature total_payments for person Index(['LAY KENNETH L', 'TOTAL'], dtype='object')
Outliers capped for feature loan_advances for person Index(['LAY KENNETH L', 'TOTAL', 'PICKERING MARK R', 'FREVERT MARK A'], dtype='object')
Outliers capped for feature bonus for person Index(['TOTAL'], dtype='object')
Outliers capped for feature restricted_stock_deferred for person Index(['TOTAL', 'BHATNAGAR SANJAY'], dtype='object')
Outliers capped for feature deferred_income for person Index(['TOTAL'], dtype='object')
Outliers capped for feature total_stock_value for person Index(['TOTAL'], dtype='object')
Outliers capped for feature expenses for person Index(['TOTAL'], dtype='object')
Outliers capped for feature from_poi_to_this_person for person Index(['LAVORATO JOHN J', 'DIETRICH JANET R'], dtype='object')
Outliers capped for feature exercised_stock_options for person Index(['TOTAL'], dtype='object')
Outliers capped for feature from_messages for person Index(['KAMINSKI WINCENTY J', 'KEAN STEVEN J'], dtype='object')
Outliers capped for feature other for person Index(['TOTAL'], dtype='object')
Outliers capped for feature from_this_person_to_poi for person Index(['BECK SALLY W', 'LAVORATO JOHN J', 'DELAINEY DAVID W', 'KEAN STEVEN J'], dtype='object')
Outliers capped for feature long_term_incentive for person Index(['TOTAL'], dtype='object')
Outliers capped for feature shared_receipt_with_poi for person Index(['LAVORATO JOHN J', 'SHAPIRO RICHARD S', 'WHALLEY LAWRENCE G', 'BELDEN TIMOTHY N'], dtype='object')
```

IPython Console History

2. What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that does not come ready-made in the dataset -- explain what feature you tried to make, and the rationale behind it. (You do not necessarily have to use it in the final analysis, only engineer and test it.) In your feature selection step, if you used an algorithm like a decision tree, please also give the feature importances of the features that you use, and if you used an automated feature selection function like SelectKBest, please report the feature scores and reasons for your choice of parameter values. [relevant rubric items: "create new features", "intelligently select features", "properly scale features"]

I decided to use all the numeric features in the POI identifier so that I could max out the classifiers with as much data as possible. I did use SelectKBest at first to determine the "best" features, and here is the result of the test, including the new feature:

```
Feature Importance Scores:
salary: 0.0703875282346141
to_messages: 1.0522029684712342
deferral_payments: 0.5078232836943501
total_payments: 2.499574642988218
loan_advances: 0.60508473015329
bonus: 0.17460771677836553
restricted_stock_deferred: 0.003744986631484842
deferred_income: 0.8559620545308682
total_stock_value: 3.560738080399393
expenses: 0.7430196598384001
from_poi_to_this_person: 5.207131180050129
exercised_stock_options: 6.6285119276760245
from_messages: 0.8087876563649865
other: 0.08612690911299588
from_this_person_to_poi: 0.28515666562628567
long_term_incentive: 0.05529990503231443
shared_receipt_with_poi: 4.727499483122773
restricted_stock: 0.742556378859582
director_fees: 0.5369391574543995
bonus_to_salary: 0.1540894179729129
```

I did create a new feature which I included in the tests – this was bonus / salary. I wanted a number that would express any bonus they made relative to their salary, as people with higher salaries might get higher bonuses.

However, this new feature was not of high relevance according to selectKBest. I decided to use the top 8 features, as well as poi as the label. Bonus_to_salary was not included due to its SelectKBest score.

After this, I did scale the features_test and features_train variables before using the machine learning algorithms on the data. The reason I scaled the data is it needs to be scaled for SVC, and in addition, I read online that it doesn't hinder the other ones like RandomForest or AdaBoost.

3. What algorithm did you end up using? What other one(s) did you try? How did model performance differ between algorithms? [relevant rubric item: “pick an algorithm”]

I tried RandomForest, SVC and AdaBoost because those seemed like good choices to use when we have a binary label. RandomForest did its job, but adaboost and SVC fell short.

Here are the evaluation metrics for them all and best grid search parameters:

```
bonus_to_salary: 0.1540894179729129
RandomForest:
Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}
Best Score: 0.1
Accuracy: 0.9545454545454546
Precision: 1.0
Recall: 0.3333333333333333
F1-score: 0.5
Create classifier
Create grid search
Fit the grid search
SVC:
Best Parameters: {'C': 100, 'kernel': 'linear'}
Best Score: 0.41428571428571426
Accuracy: 0.7954545454545454
Precision: 0.0
Recall: 0.0
F1-score: 0.0
AdaBoost:
Best Parameters: {'learning_rate': 1.0, 'n_estimators': 100}
Best Score: 0.2831746031746032
Accuracy: 0.8636363636363636
Precision: 0.2
Recall: 0.3333333333333333
F1-score: 0.25
```

4. What does it mean to tune the parameters of an algorithm, and what can happen if you don't do this well? How did you tune the parameters of your particular algorithm? What parameters did you tune? (Some algorithms do not have parameters that you need to tune -- if this is the case for the one you picked, identify and briefly explain how you would have done it for the model that was not your final choice or a different model that does utilize parameter tuning, e.g. a decision tree classifier). [relevant rubric items: “discuss parameter tuning”, “tune the algorithm”]

The parameters used when running an algorithm can create vastly different results depending on what is chosen. This means that the algorithm can perform poorly if it is not optimized. Therefore, I used a grid search on all three algorithms in order to determine the optimal parameters.

ParamGrid for RandomForest:

```
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
```

```
'min_samples_leaf': [1, 2, 4]
}
```

ParamGrid for SVC:

```
param_grid = {'C': [0.1, 1, 10, 100],
              'kernel': ['linear', 'rbf']}
```

ParamGrid for AdaBoost:

```
param_grid = {'n_estimators':[50, 100, 200],
              'learning_rate':[0.1, 0.5, 1.0]}
```

5. What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis? [relevant rubric items: "discuss validation", "validation strategy"]

I validated all three chosen algorithms by dividing the data into test and training data, and then using the test results to determine various evaluation metrics for each algorithm, including accuracy, recall, precision and f1 score. Since the algorithms were trained on the training data, we can test them on the test data and validate that way.

If you do not use this method, you have not tested the algorithm on data it has not seen yet, only on the data it was trained on. This can make the algorithm over-perform.

If you don't look at precision or recall, you might have a high accuracy that is irrelevant. In the context of this project, for example, if you said no one was a poi, you would have a high accuracy just because there are so many non-pois.

The `test_classifier()` function in `tester.py` uses a stratified shuffle split in order to cross-validate the data.

6. Give at least 2 evaluation metrics and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance. [relevant rubric item: "usage of evaluation metrics"]

The best performance was with RandomForest using optimal parameters and it was the following:

I will be using accuracy, precision, recall and f1 scores as evaluation metrics.

RandomForest:

Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}

Best Score: 0.1

Accuracy: 0.9545454545454546

Precision: 1.0

Recall: 0.3333333333333333

This was our star algorithm.

The accuracy of .95 means that the algorithm was able to make the correct determination 95% of the time.

The precision of 1.0 means that 100% of the time the algorithm chose someone as a person of interest, they were one.

The recall of .33 means that 33% of the time, the people of interest were discovered by the algorithm.

The f1 score is the harmonic mean of recall and precision. This is used to balance the trade-off between recall and precision by using one of their means.

Notes on the code:

The first thing that I do is transform the dictionary into a dataframe, which I found easier to use when writing outlierCleaner and replace_nan_with_mean, both of which are included in the outlier_cleaner.py file.

outlierCleaner uses a z-score method to cap any outliers above a z-score of 3 to a value at a z-score of 3.

Replace_nan_with_mean calculates the mean of each column and replaces any nan values with that.

Next, selectKBest is used in order to do feature analysis and those scores are printed out.

After this, I scale the features_test and features_train variables

Next, I try three different classifiers – RandomForest, adaboost and SVC. I used search grids on each one to find the best values to use for the parameters.

NAN VALUES were present in the following columns and were replaced by the column mean:

```
Function activated
NAN values replaced for salary
NAN values replaced for to_messages
NAN values replaced for deferral_payments
NAN values replaced for total_payments
NAN values replaced for loan_advances
NAN values replaced for bonus
NAN values replaced for restricted_stock_deferred
NAN values replaced for deferred_income
NAN values replaced for total_stock_value
NAN values replaced for expenses
NAN values replaced for from_poi_to_this_person
NAN values replaced for exercised_stock_options
NAN values replaced for from_messages
NAN values replaced for other
NAN values replaced for from_this_person_to_poi
NAN values replaced for long_term_incentive
NAN values replaced for shared_receipt_with_poi
NAN values replaced for restricted_stock
NAN values replaced for director_fees
poi      salary  ...  restricted_stock  director_fees
METTS MARK      False  365788.000000  ...      585062.0  166804.882353
BAXTER JOHN C   False  267102.000000  ...      3942714.0  166804.882353
ELLIOTT STEVEN  False  170941.000000  ...      1788391.0  166804.882353
CORDES WILLIAM R False  562194.294737  ...      386335.0  166804.882353
```

Here were the results of the other classifiers:

SVC:

Best Parameters: {'C': 100, 'kernel': 'linear'}

Best Score: 0.41428571428571426

Accuracy: 0.7954545454545454

Precision: 0.0

Recall: 0.0

F1-score: 0.0

AdaBoost:

Best Parameters: {'learning_rate': 1.0, 'n_estimators': 100}

Best Score: 0.2831746031746032

Accuracy: 0.8636363636363636

Precision: 0.2

Recall: 0.3333333333333333

F1-score: 0.25